

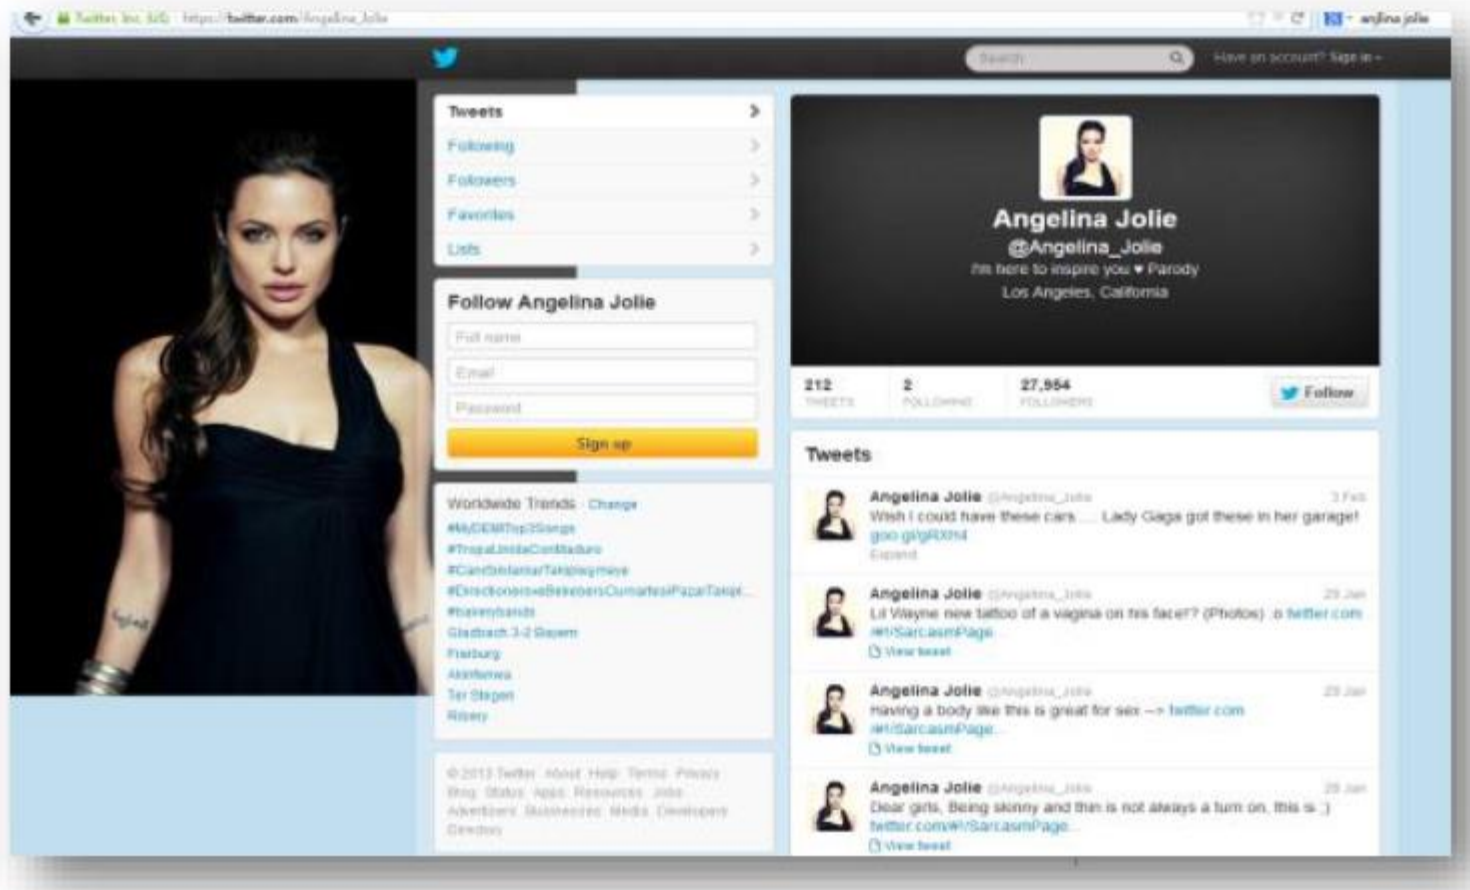
HBase

# Agenda

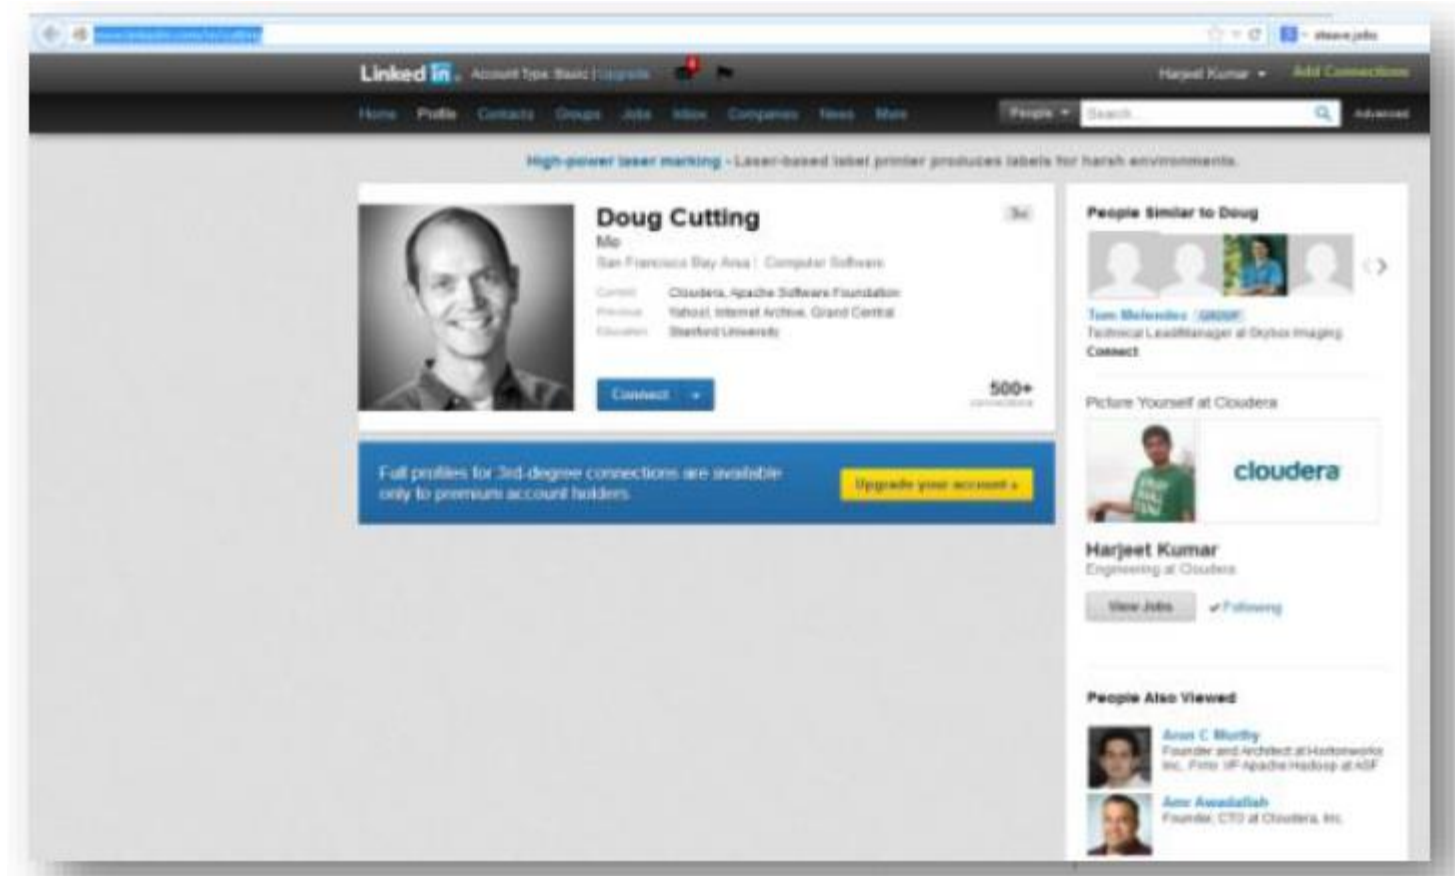
---

- ✓ Motivation for HBase
  - What and Why
- ✓ Hbase
  - Overview
  - Architecture
  - Features
  - HBase Commands
  - Data Loading Techniques
- ✓ Zookeeper

# Problems in Real World



# Problems in Real World

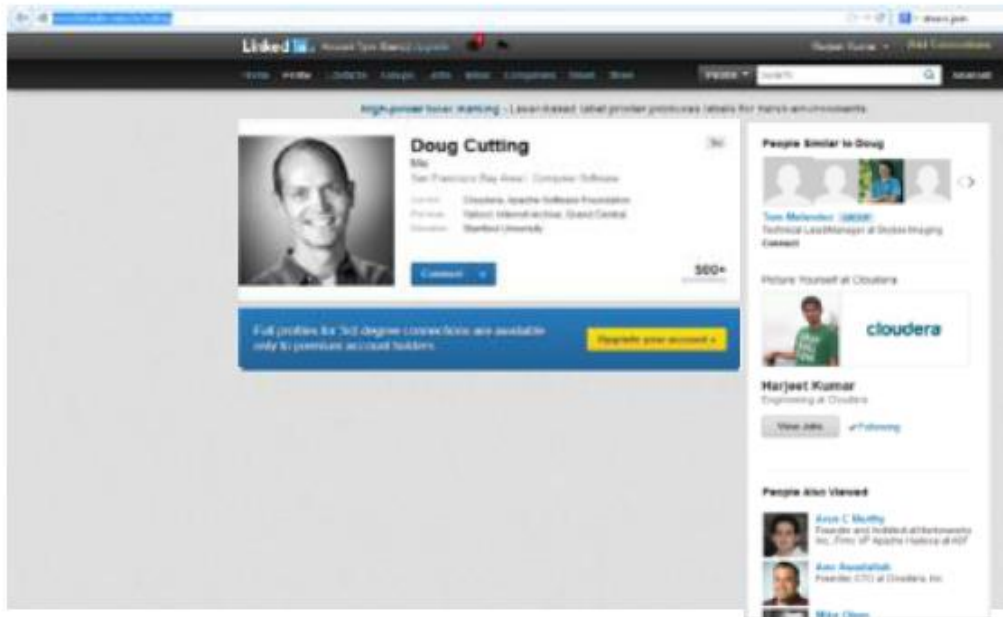


# So, What Is Common?

---

- ✓ Huge Data
- ✓ Fast Random access
- ✓ Structured Data
- ✓ Variable Schema
- ✓ Need of Compression
- ✓ Need of Distribution (Sharding)

# How RDBMS will solve this?



## Users

Id

Name

Sex

age

## Connections

User\_id

Connection\_id

type

# Characteristics Of Probable Solution

---

- ✓ Distributed Database
- ✓ Sorted Data
- ✓ Sparse Data Store
- ✓ Automatic sharding

# HBase Definition

---

✓ **HBase is a key/value store. Specifically it is:**

Sparse

Distributed

Multi-  
dimensional

Sorted  
Map

Consistent



# Hbase v/s RDBMS

HBase	RDBMS
Column-oriented	Row oriented (mostly)
Flexible schema, add columns on the fly	Fixed schema.
Good with sparse tables,	Not optimized for sparse tables.
Joins using MR –not optimized	Optimized for joins.
Tight integration with MR	Not really...
Horizontal scalability –just add hardware	Hard to shard and scale
Good for semi-structured data as well as structured data	Good for structured data

# HBase—The Map Is Multi-Dimensional

```
{  
  "1" : {  
    "A" : "x",  
    "B" : "z"  
  },  
  "aaaaa" : {  
    "A" : "y",  
    "B" : "w"  
  },  
  "aaaab" : {  
    "A" : "world",  
    "B" : "ocean"  
  },  
  "xyz" : {  
    "A" : "hello",  
    "B" : "there"  
  },  
}
```

Top Level Key/map pair is called row

A and B are called Column families

# Row vs Column Oriented DBs

URLS					
url-id INTEGER PK	url VARCHAR(4096)	ref_short_id CHAR(8)	title VARCHAR(200)	description VARCHAR(400)	content TEXT
1	http://hbase.apache.org	3fG4J	HBase Home	Great tool!	<html><head><title>Hbas e Home</ti...
2	http://larsgeorge.com	1337	Lineland	<NULL>	<html><body>Newest Posts...
3	http://foobar.com/index.html	Hf34h	<NULL>	Read about it...	404 Page not found.
4	http://cnn.com/page123.html	0o001	Sport News	Soccer News	<html><body>Results, Reviews,...

SQL Schema

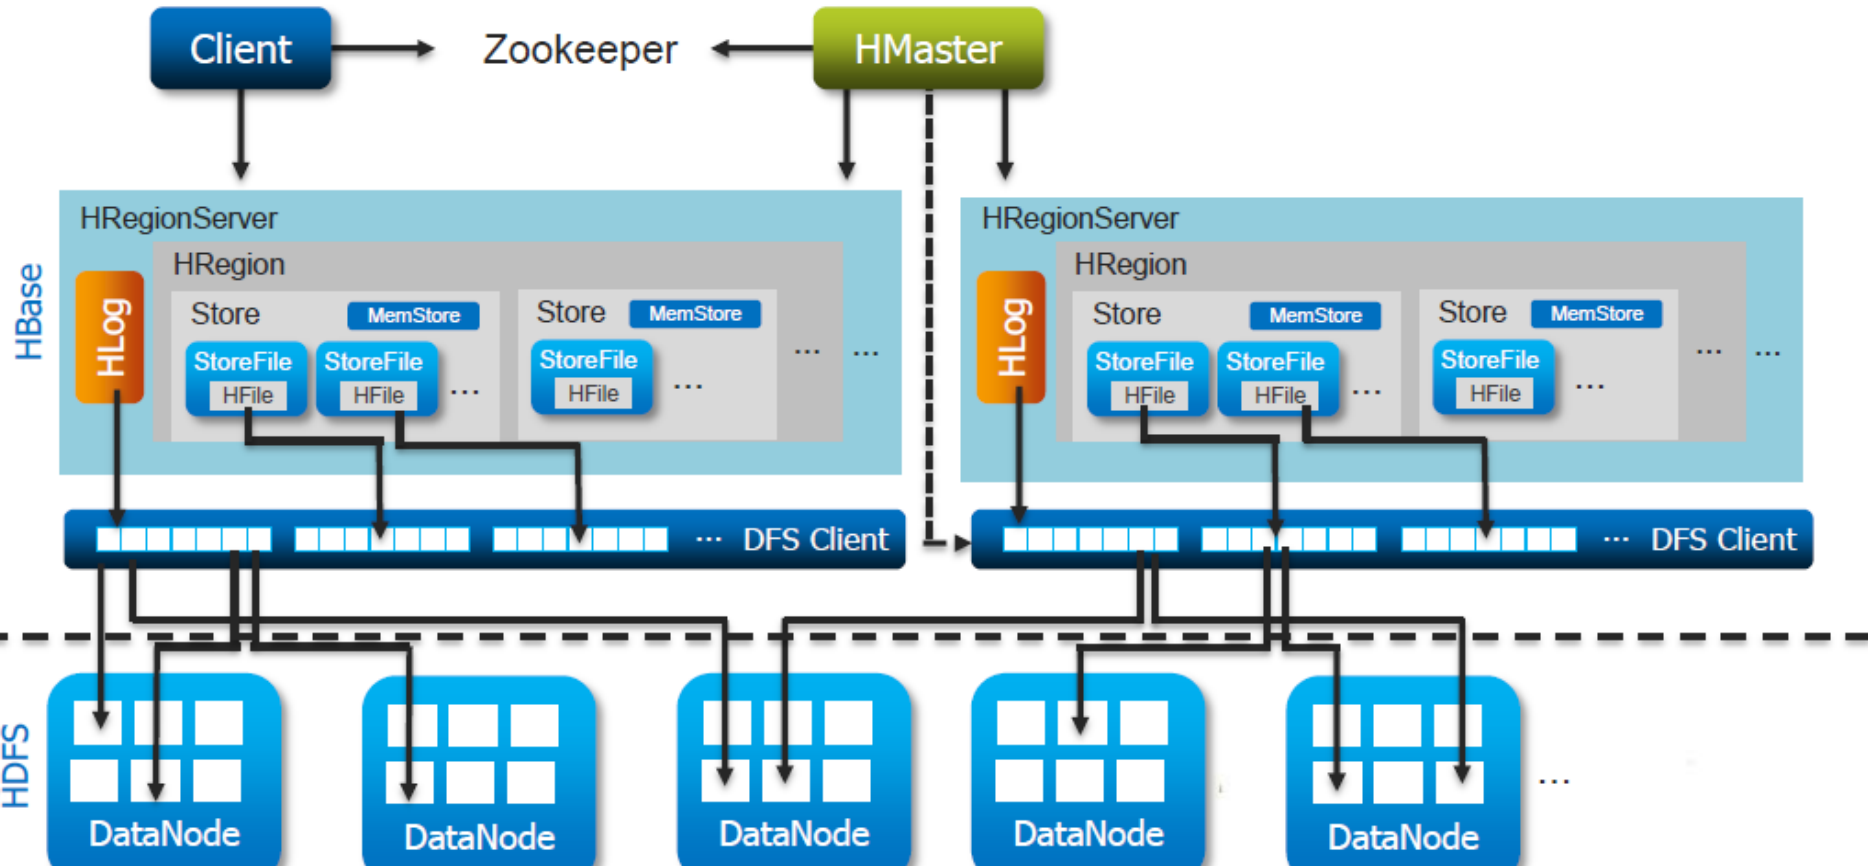
Column Oriented Storage

Col 1: url	<a href="http://hbase.apache.org">http://hbase.apache.org</a>	<a href="http://larsgeorge.com">http://larsgeorge.com</a>	<a href="http://foobar.com/index.html">http://foobar.com/index.html</a>	http://cnn.com/page12...	...
Col 2: ref_short_id	3fG4J	1337	Hf34h	0o001	...
Col 3: title	HBase Home	Lineland	<NULL>	Sport News	...
Col 4: description	Great tool!	<NULL>	Read about it...	Soccer News	...
Col 5: content	<html><head><title>HBA...	<<html><body>Newest Po...	404 Page not found.	<html><body>Results,...	...

# Data Model

Row key	Personal_data		demographic	
Persons ID	Name	Address	Birth Date	Gender
1	H. Houdini	Budapest	1926-10-31	M
2	D. Copper		1956-09-16	M
3	Merlin		1136-12-03	F
4	.....	.....	.....	M
500,000,000	F. Cadillac	Nevada	1964-01-07	M

# HBase Storage Architecture



# Loading Data Into HBASE Using Pig



## Example:

input.csv

```
1, fname1, lname1  
2, fname2, lname2  
3, fname3, lname3
```

```
raw_data = LOAD 'input.csv' USING PigStorage( ',' ) AS (  
    listing_id: chararray,  
    fname: chararray,  
    lname: chararray);  
dump raw_data;  
STORE raw_data INTO 'hbase://sample_names' USING  
org.apache.pig.backend.hadoop.hbase.HBaseStorage('info:  
fname info:lname');
```

# Loading Data Into HBASE Using Sqoop



**Sqoop can be used to directly import data from RDBMS to Hbase:**

## **Example:**

```
sqoop import
--connect jdbc:mysql://<ip address>\<database name>
--username <username_for_mysql_user> --password
<Password>
--table <mysql_table name>
--hbase-table<hbase_target_table_name>
--column-family <column_family_name>
--hbase-row-key <row_key_column>
--hbase-create-table
```

# Problem

---



## Problem:

In a distributed environment, getting processes to act in any kind of synchrony is an extremely hard problem.

**For example,** simply having a set of processes wait until they've all reached the same point in their execution – a kind of distributed barrier– is surprisingly difficult to do correctly.



## Solution:

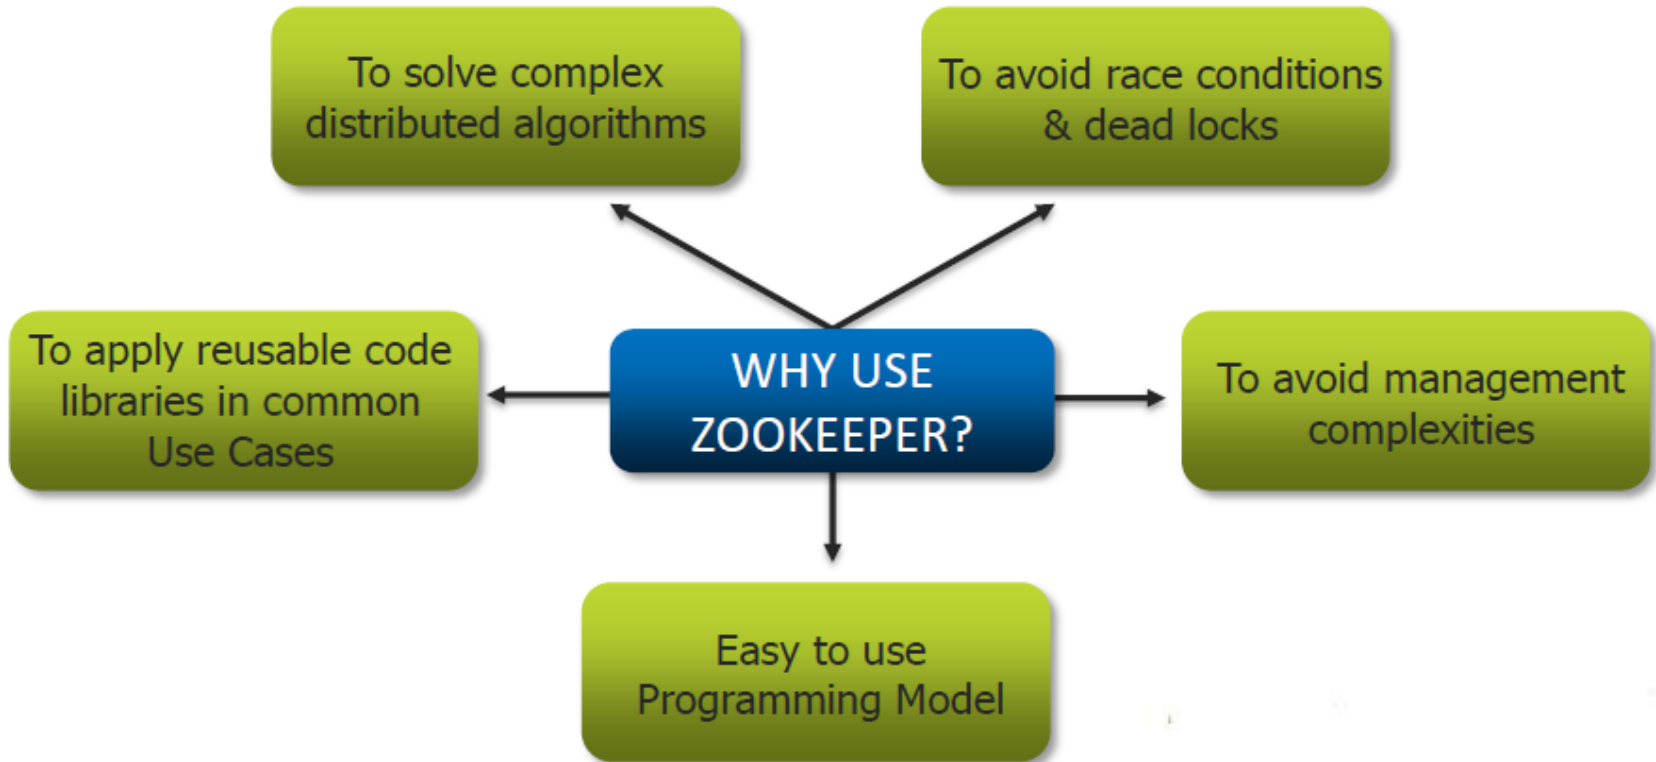
**ZooKeeper** offers an API to facilitate this sort of distributed coordination.

**For example,** it is often used to serve locks to client processes – locks are just another kind of coordination primitive.



# Why Use Zookeeper?

---



# What Is A ZooKeeper?

## A Reliable, Scalable Distributed Coordination System

ZooKeeper is a sort of central nervous system for distributed systems where the role of the brain is played by the [coordination service](#), axons are the [network](#), [processes](#) are the monitored and controlled body parts, and [events](#) are the hormones and neurotransmitters used for messaging.

***Apache ZooKeeper*** is a software project of the Apache Software Foundation, providing an open source distributed configuration service, synchronization service and naming registry for large distributed systems.

Every complex distributed application needs a coordination and orchestration system of some sort, so the ZooKeeper folks at Yahoo decide to build a good one and open source it for everyone to use!



# Distribute Coordination System

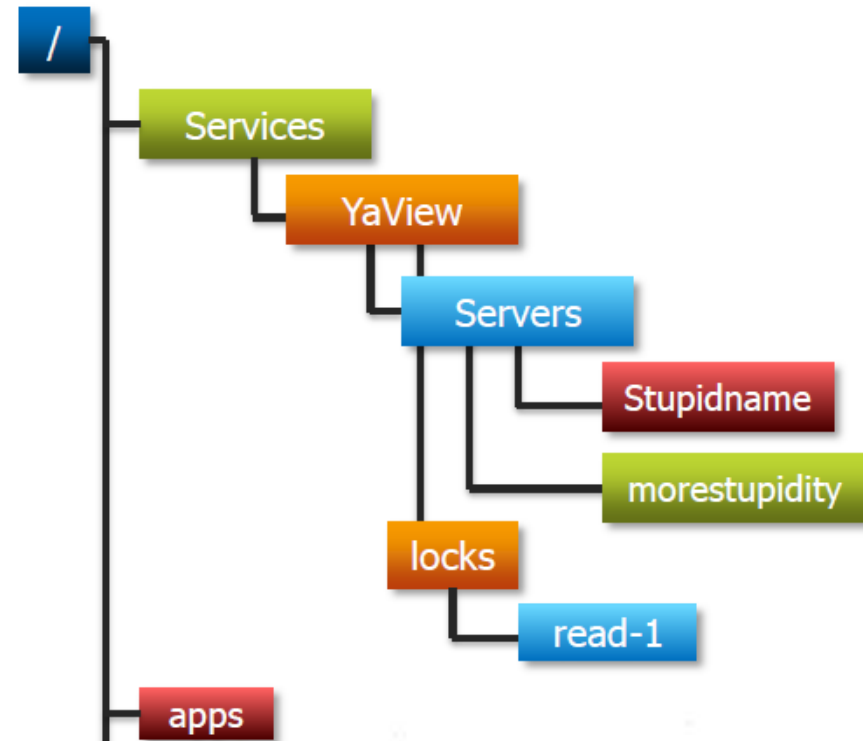
---

## ✓ Why would you ever need a Distribute Coordination System?

Because you can't get these guarantees from an *event system* plopped on top of a database and these are the sort of guarantees you need in a *complex distributed system* where connections drop, nodes fail, retransmits happen, and chaos rules the day.

# ZooKeeperData Model

- ✓ Hierarchal namespace (like a File System)
- ✓ Each znode has data and children
- ✓ Data is read and written in its entity



# Hbase Commands

```
hbase(main):018:0> create 'student', {NAME => 'studinfo',VERSIONS => 5},{NAME => 'studaddress'}
0 row(s) in 1.5880 seconds
```

```
hbase(main):008:0> put 'employee','1000','empinfo:firstname','Ankit'
0 row(s) in 0.0980 seconds
```

```
hbase(main):015:0> delete 'employee','1000','empinfo:firstname'
0 row(s) in 0.0070 seconds
```

```
hbase(main):032:0> get 'student','stu0001','studinfo:salary'
COLUMN                                CELL
studinfo:salary                       timestamp=1391881336240, value=2000
1 row(s) in 0.0360 seconds
```

```
hbase(main):038:0> get 'student','stu0001',{COLUMN => 'studinfo:salary', VERSIONS => 3}
COLUMN                                CELL
studinfo:salary                       timestamp=1391881524160, value=4000
studinfo:salary                       timestamp=1391881475201, value=3000
studinfo:salary                       timestamp=1391881336240, value=2000
3 row(s) in 0.0620 seconds
```

```
hbase(main):048:0> count 'student'
1 row(s) in 0.0520 seconds
```

# Hbase Commands

```
hbase(main):049:0> exists 'student'
Table student does exist
0 row(s) in 0.1180 seconds
```

```
hbase(main):051:0> status 'detailed'
version 0.90.6-cdh3u6
0 regionsInTransition
1 live servers
  cloudera-vm:60020 1391877650590
    requests=0, regions=4, usedHeap=24, maxHeap=998
    .META.,,1
      stores=1, storefiles=0, storefileSizeMB=0, memstoreSizeMB=0, storefileIndexSizeMB=0
employee,,1391877808132.225413b01230b8e959269742d5305819.
  stores=2, storefiles=2, storefileSizeMB=0, memstoreSizeMB=0, storefileIndexSizeMB=0
-ROOT-,0
  stores=1, storefiles=1, storefileSizeMB=0, memstoreSizeMB=0, storefileIndexSizeMB=0
student,,1391881252322.6a322d97031d67b211a51db13cbfdd60.
  stores=2, storefiles=0, storefileSizeMB=0, memstoreSizeMB=0, storefileIndexSizeMB=0
0 dead servers
```

```
hbase(main):052:0> truncate 'employee'
Truncating 'employee' table (it may take a while):
- Disabling table...
- Dropping table...
- Creating table...
0 row(s) in 4.0610 seconds
```

# Hbase Commands

```
hbase(main):054:0> drop 'employee'

ERROR: Table employee is enabled. Disable it first.'

Here is some help for this command:
Drop the named table. Table must first be disabled. If table has
more than one region, run a major compaction on .META.:

hbase> major_compact ".META."

hbase(main):055:0> disable 'employee'
0 row(s) in 2.0760 seconds

hbase(main):056:0> enable 'employee'
0 row(s) in 2.0760 seconds
```

```
alter 't1',{NAME=>'cf1',VERSIONS=>5}
```

```
alter 't1',{NAME=>'cf1',METHOD=>'delete'}
```

# Zookeeper

---

- Hbase zkcli
- Ls /
- Node creation : Create /jaymin <value> and then get /jaymin
- Get /hbase/root-region-server
- Ls /hbase/rs – to get all region servers
- Types of znodes – ephemeral (when rs dies, this node gets deleted from zookeeper) and persistent